

Computer Intensive Methods using R

Part 4: bootstrap tests

Prof. Dr. Ziv Shkedy

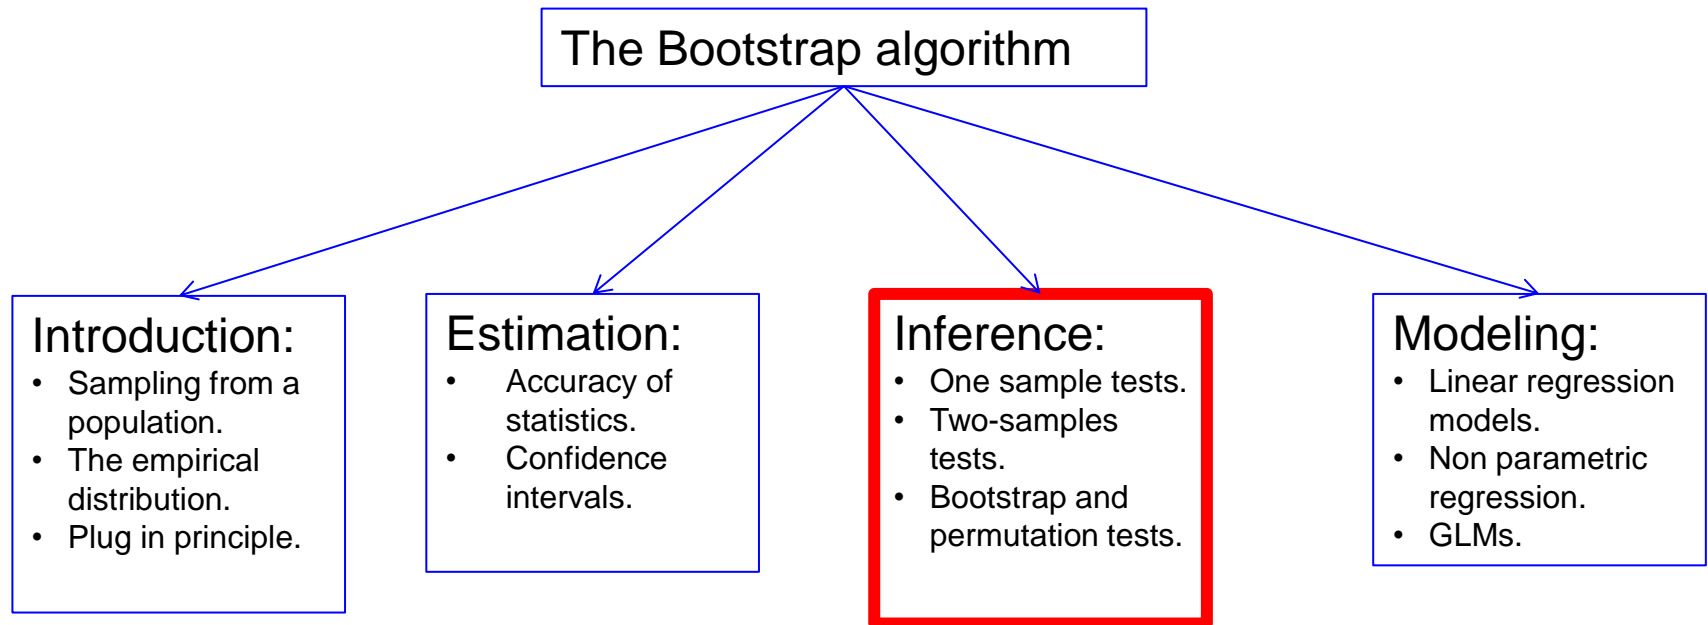
Master of Statistics
Hasselt University

General Information

Overview of the course


- Bootstrap tests.
 - One sample problem.
 - Two samples problem.
- Permutation tests:
 - Two samples problem.
- Correlation test.
- Binary data.

Overview of the course (part 1)



Reference

- Bradley Efron and Robert J. Tibshirani (1994): An introduction to bootstrap.
- Davison A.C. and Hinkley D.V: Bootstrap Methods and Their Application.




Course materials

- Slides.
- R program.
- R datasets & External datasets.
- YouTube tutorials.
- Videos for the classes (highlights of each class in the course).

YouTube tutorials

- YouTube tutorials about bootstrap using R:
 1. One-sample bootstrap CI for the mean (host: [LawrenceStats](#)): <https://www.youtube.com/watch?v=ZkCDYAC2iFg>.
 2. Using the non-parametric bootstrap for regression models in R (host: [lan Dworkin](#)): <https://www.youtube.com/watch?v=ydtOTctg5So>.
 3. Performing the Non-parametric Bootstrap for statistical inference using R (host: [lan Dworkin](#)): <https://www.youtube.com/watch?v=TP6r5CTd9yM>
 4. Using the sample function in R for resampling of data - absolute basics (host: [lan Dworkin](#)): <https://www.youtube.com/watch?v=xE3KGVt6VLE>
 5. Permutation tests in R - the basics (host: [lan Dworkin](#)): <https://www.youtube.com/watch?v=ZiQdzWB12Pk>.
 6. Bootstrap Sample Technique in R software (host: [Sarveshwar Inani](#)): <https://www.youtube.com/watch?v=tb6wb9ZdPH0>
 7. Bootstrap confidence intervals for a single proportion (host: [LawrenceStats](#)): <https://www.youtube.com/watch?v=ubX4QEPqx5o>
 8. Bootstrapped prediction intervals (host: [James Scott](#)): https://www.youtube.com/watch?v=c3gD_PwsCGM.
- <https://www.youtube.com/watch?v=gcPlyeqymOU>

Bootstrap tests



Topics

- Testing of hypothesis in one sample.
- Testing equality of the means in two populations: algorithm 16.1.
- Testing equality of the means in two populations: algorithm 16.2.
- Correlation test (the cars data).
- Ratio test and pair t-test ([the city population data](#)).
- Hypothesis testing for binary data (two populations).

The
mouse
data.

External dataset.

Bootstrap tests

As before the bootstrap mechanism should reflect the original data generation mechanism and, in this setting of testing hypotheses, there is a second important rule

The bootstrap simulation should satisfy the null hypothesis.


Bootstrap tests

The idea is to generate (repeatedly) new bootstrap data, reflecting the null hypothesis, recalculate the test statistic and in this way to simulate the null distribution of the test statistic.

These bootstrap test values can then be used to compute a p-value.

Example 1: One-sample problem

The one sample problem



Population



sample

X_1, X_2, \dots, X_n

Distribution of X in the population

$$X_i \sim G(\theta)$$

$$E(X_i) = \mu$$

We wish to test
the hypotheses

$$H_0 : \mu = \mu_0$$

$$H_1 : \mu > \mu_0$$



One Sample problem

$$H_0 : \mu = \mu_0$$

$$H_1 : \mu > \mu_0$$



Observed sample

$$x_1, x_2, \dots, x_n$$

Test statistic

$$t = \frac{\bar{x} - \mu_0}{se(\bar{x})}$$




Distribution of the test statistic under **the null hypothesis**

$$t \sim t_{(n-1)}$$

One sample problem: inference

Rejection region




P value

$$P_{H_0}(t > t^{obs}) < \alpha$$


The p value

Distribution of the test statistic under the null



Bootstrap test

Distribution of the test statistic under the null



We would like to approximate the distribution of the test statistic under the null hypothesis using bootstrap.

The bootstrap simulation should satisfy the null hypothesis.

Example: survival time of mice

- Z: survival time of 7 mice, in days, after surgery are:

```
> z  
[1] 94 197 16 38 99 141 23
```

```
> mean(z)  
[1] 86.85714
```

```
> sd(z)  
[1] 66.76683
```

One Sample problem (two sided test)

We wish to test the null hypothesis:

$$H_0 : \mu = 129$$

$$H_1 : \mu \neq 129$$

Observed data:



$$z_1, z_2, \dots, z_7$$

Test statistics and the distribution under the null:

$$t = \frac{\bar{z} - \mu_0}{se(\bar{z})}$$

$$t \sim t_{(6)}$$

One sample t test

```
> z <- c(94, 197, 16, 38, 99, 141, 23)
> t.test(z,mu=129)
```

One Sample t-test

```
data:  z
t = -1.67, df = 6, p-value = 0.1460
alternative hypothesis: true mean is not
equal to 129
95 percent confidence interval:
 25.10812 148.60616
sample estimates:
mean of x
 86.85714
```


Two sided test

$$t = \frac{\bar{z} - \mu_0}{se(\bar{z})}$$

Under H0, the distribution of the test statistics:

$$t \sim t_{(6)}$$

Distribution of the test statistic under the null



```
x<-seq(from=-3,to=3,length=1000)
> dx<-dt(x,6)
> plot(x,dx,type="l")
> lines(c(-1.67,-1.67),c(0,1),lwd=3,col=2)
> pt(-1.67,6)
[1] 0.07297866
```

```
>
t.test(z,mu=129,alternative="less")
```

One Sample t-test

```
data: z
t = -1.67, df = 6, p-value = 0.07298
alternative hypothesis: true mean is
less than 129
95 percent confidence interval:
    -Inf 135.8942
sample estimates:
mean of x
86.85714
```

One sided test

Bootstrap test (1)

How can we generate the bootstrap sample **under the null hypothesis** ?

$$H_0 : \mu = 129$$

$$H_1 : \mu \neq 129$$




Assuming a normal population, the distribution of the the survival time and sample mean under the null hypothesis is given by

$$Z_i \sim N(129, \sigma^2)$$

$$\bar{Z}_i \sim N(129, \frac{\sigma^2}{7})$$


The empirical distribution under the null hypothesis



Bootstrap test (1)

How can we generate the bootstrap sample **under the null hypothesis** ?

$$H_0 : \mu = 129$$

$$H_1 : \mu \neq 129$$


We define a new variable for which the mean is equal to 129

$$\tilde{Z}_i = Z_i - \bar{Z} + \mu_0$$

$$E(\tilde{Z}_i) = \mu_0$$

```
> mz <- mean(z)
> mz
[1] 86.85714
> n <- length(z)
> z.tilde <- z - mz + 129
> mean(z.tilde)
[1] 129
```


Two sided non parametric bootstrap t-test

Algorithm for the one sample problem

Simulate B samples of size n
with replacement from \tilde{z}

Calculate the test statistic values

$$t(z^*)_1, t(z^*)_2, \dots, t(z^*)_B$$

Compute the Monte Carlo P-value

$$P = \frac{1 + \# \left\{ |\hat{t}^*| > |t^{obs}| \right\}}{B + 1}$$

```
nz<-length(z)
t.obs<-t.test(z,mu=129)$statistic
B<-5000
t.boot<-c(1:B)
for(b in 1:B)
{
  z.b<-sample(z.tilde,size=nz,replace=T)
  t.boot[b]<-
    t.test(z.b,mu=129)$statistic
}
```

For two sided test

Bootstrap p value (two sided)


$$P = \frac{1 + \#\left\{|\hat{t}^*| > |t^{obs}|\right\}}{B + 1}$$

```
> Pmc<-(1+sum(abs(t.boot) > abs(t.obs)))/(B+1)
> Pmc
[1] 0.1519696
```

```
> t.test(z,mu=129)
```

One Sample t-test

```
data: z
t=-1.67, df=6, p-value=0.1460
```



Two sided parametric bootstrap t-test

Algorithm for the one sample problem

Simulate B samples of size n
from

$$F(\theta_0)$$

Calculate the test statistic values

$$t(z^*)_1, t(z^*)_2, \dots, t(z^*)_B$$

Compute the Monte Carlo P-value

$$P = \frac{1 + \# \left\{ |\hat{t}^*| > |t^{obs}| \right\}}{B + 1}$$

For two sided test

```
Theta.0<-129
Sig.0<-ssrt(var(z))
nz<-length(z)
t.obs<-t.test(z,mu=129)$statistic
B<-5000
t.boot<-c(1:B)
for(b in 1:B)
{
  z.b<-rnorm(nz,theta.0,sig.0)  $N(\mu_0, \hat{\sigma}^2)$ 
  t.boot[b]<-t.test(z.b,mu=129)$statistic
}
```

Example 2: two-sample problem

The two samples problem


- We would like to compare the distribution of a random variable in two populations (for example: EEG for two treatments in a clinical trial).
- Our main interest is to test if the distribution of the variable of primary interest is the same in the two populations.

Two populations – two distribution functions

$$H_0 : F = G$$

$$H_1 : F > G$$

We assume that F is a
shift of the distribution
G



Two populations – two distribution functions


$$H_0 : F = G$$

$$H_1 : F > G$$

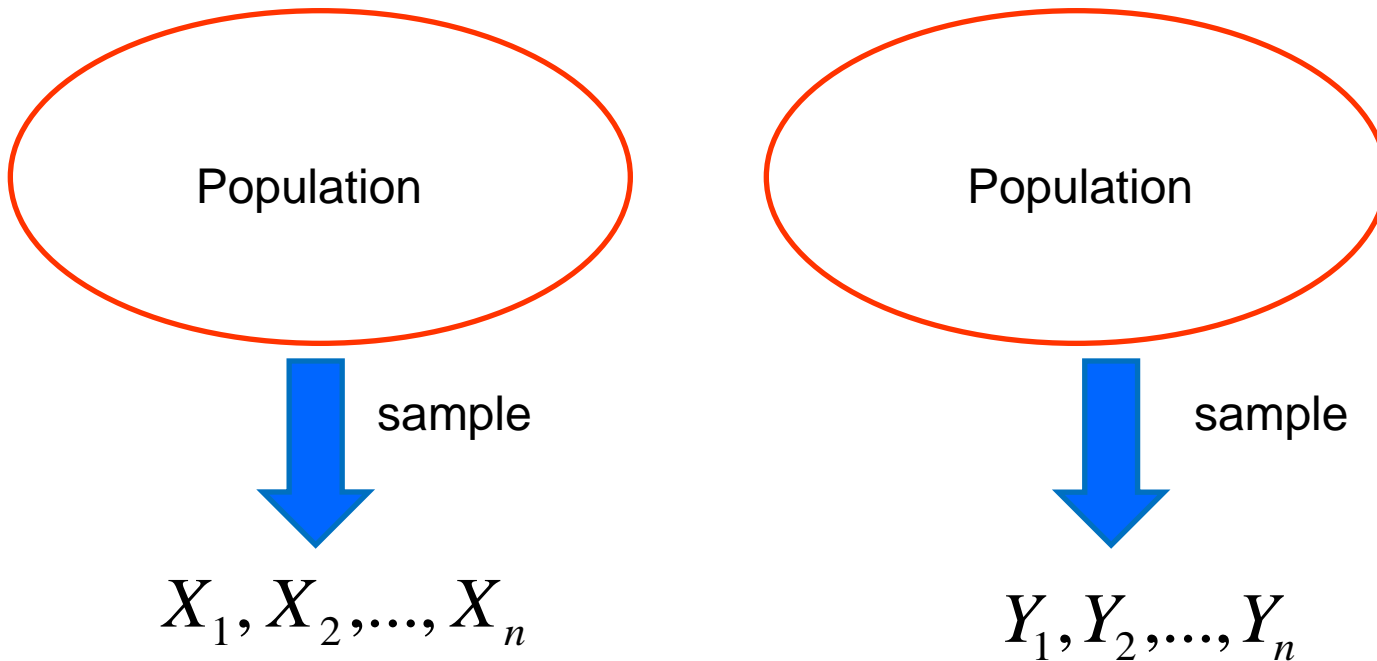


$$H_0 : \mu_1 = \mu_2$$

$$H_1 : \mu_1 > \mu_2$$



Populations and samples




$$(Z_1, Z_2, \dots, Z_n, Z_{n+1}, \dots, Z_{n+m}) = (X_1, X_2, \dots, X_n, Y_1, \dots, Y_m)$$

↑

The mouse data

- A randomized experiment with 16 mouse:
 - 7 to treatment group
 - 9 to control group.
- Treatment was intended to prolong survival after a test surgery.



```
> y <- c(10, 27, 31, 40, 46, 50, 52, 104, 146)
> x <- c(16, 23, 38, 94, 99, 141, 197)
> mean(y)
[1] 56.22222
> mean(x)
[1] 86.85714
```

Two samples problem

$$H_0 : \mu_1 = \mu_2$$

$$H_1 : \mu_1 > \mu_2$$

$$t = \frac{\bar{x} - \bar{y}}{\hat{\sigma} \sqrt{\frac{1}{n} + \frac{1}{m}}}$$

$$t \sim t_{(n+m-2)}$$

```
> y
[1] 10 27 31 40 46
50 52 104 146
> x
[1] 16 23 38 94 99
141 197

> mean(y)
[1] 56.22222
> mean(x)
[1] 86.85714
```

Two-sample t test

```
> t.test(x,y,alternative="greater",var.equal=TRUE)
```

Two Sample t-test

data: x and y

t = 1.1214, df = 14, p-value = 0.1405

alternative hypothesis: true difference in means is greater than
0

95 percent confidence interval:

-17.48178 Inf


sample estimates:

mean of x mean of y

86.85714 56.22222

We do not reject the null hypothesis

Bootstrap algorithm 1



Algorithm 16.1

Algorithm 16.1



$$z = (y, x)$$

Simulate B samples of size $m + n$ with replacement from z .

The first n are considered as observations x^* and the remaining m as observations y^* .

Calculate the test statistic for each bootstrap sample

$$t(z^*)_1, t(z^*)_2, \dots, t(z^*)_B$$

Two samples problem: test statistics I

$$H_0 : \mu_1 = \mu_2$$

$$H_1 : \mu_1 > \mu_2$$

Test statistic

$$t = \bar{x} - \bar{y}$$

```
> t.obs <- mean(x) - mean(y)
> t.obs
[1] 30.63492
```


Simulate B samples of size $m+n$ with replacement from z .

The first n are considered as observations x^* and the remaining m as observations y^* .

Calculate the test statistic for each bootstrap sample

$$t(z^*)_1, t(z^*)_2, \dots, t(z^*)_B$$

Two-sample bootstrap t test (one sided)



```
> (sum(t.boot>t.obs)+1)/(B+1)
[1] 0.09318136
```

```
z <- c(x, y)
m <- length(x)
n <- length(y)
mn <- m + n
t.obs <- mean(x) - mean(y)
t.obs
B<-5000
t.boot<-c(1:B)
for(b in 1:B)
{
  z.b<-sample(z,size=mn,replace=T)
  x.b<-z.b[1:n]
  y.b<-z.b[(n+1):mn]
  t.boot[b]<-mean(x.b)-mean(y.b)
}
hist(t.boot,nclass=50,probability=T)
lines(c(t.obs,t.obs),c(0,1),lwd=3,col=2)
```

Two samples problem: test statistic II

$$H_0 : \mu_1 = \mu_2$$

$$H_1 : \mu_1 > \mu_2$$

What is the problem with the previous test statistic ?

```
> t.test(x,y,alternative="greater")
```

$$t = \frac{\bar{x} - \bar{y}}{\hat{\sigma} \sqrt{\frac{1}{n} + \frac{1}{m}}}$$

$$t \sim t_{(n+m-2)}$$




Assuming
equal variance

Welch Two Sample t-test

```
data:  x and y
t = 1.0591, df = 9.645, p-value = 0.1577
alternative hypothesis: true difference
in means is greater than 0
95 percent confidence interval:
 -21.98945      Inf
sample estimates:
mean of x mean of y
 86.85714  56.22222
```


Two-samples bootstrap t test (one sided)




```
> B<-5000
> t.boot<-c(1:B)
> for(b in 1:B)
+ {
+   z.b<-sample(z,size=mn,replace=T)
+   x.b<-z.b[1:n]
+   y.b<-z.b[(n+1):mn] ←
+   t.boot[b]<-t.test(x.b,y.b)$statistic
+ }
> hist(t.boot,nclass=50,probability=T)
> lines(c(t.obs,t.obs),c(0,1),lwd=3,col=2)
```

```
> Pmc<-(1+sum(t.boot > t.obs))/(B+1)
> Pmc
[1] 0.1445711
```


How many bootstrap samples ?

Histogram of t.boot



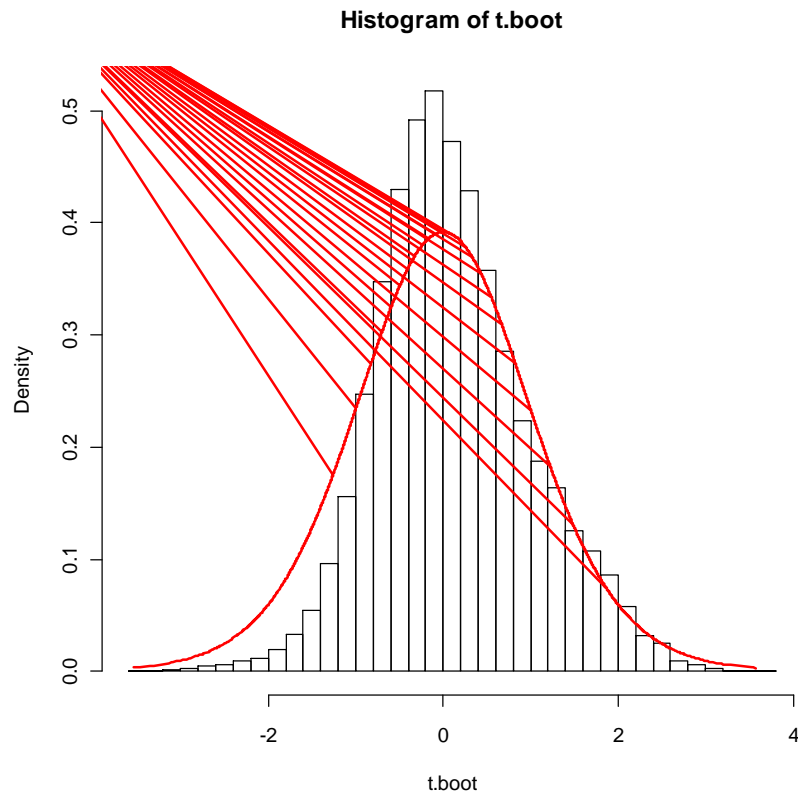
B=15000

Histogram of t.boot



B=150

B=25000



Two Sample t-test

data: x and y

t = 1.1214, df = 14, p-value = 0.1405

alternative hypothesis: true

difference in means is greater than
0

95 percent confidence interval:

-17.48178 Inf

sample estimates:

mean of x mean of y


86.85714 56.22222

```
> Pmc<-(1+sum(t.boot > t.obs))/(B+1)
```

```
> Pmc
```

```
[1] 0.1484341
```

Bootstrap algorithm 2



Algorithm 16.2

Two samples problem: test statistics I

$$z = (y, x)$$

$$\bar{z} = \frac{1}{n} \sum_{i=1}^n z_i$$

$$\tilde{x}_i = x_i - \bar{x} + \bar{z} \quad \leftarrow$$

$$\tilde{y}_i = y_i - \bar{y} + \bar{z} \quad \leftarrow$$

Simulate B samples with replacement of size **n** from **x-tilde** and size **m** from **y-tilde**.

Calculate the test statistic for each bootstrap sample:

$$t(z^*)_1, t(z^*)_2, \dots, t(z^*)_B$$

$$t^* = \frac{\bar{x}^* - \bar{y}^*}{\hat{\sigma}^* \sqrt{\frac{1}{n} + \frac{1}{m}}}$$

Non parametric bootstrap (algorithm 16.2)


```
> y <- c(10, 27, 31, 40, 46, 50, 52, 104, 146)
> x <- c(16, 23, 38, 94, 99, 141, 197)
> z<-c(x,y)
> m <- length(x)
> n <- length(y)
> my<-mean(y)
> mx<-mean(x)
> mz<-mean(z)
> x.tilde<-x-mx+mz
> y.tilde<-y-my+mz
> mean(x.tilde)
[1] 69.625
> mean(y.tilde)
[1] 69.625
```

$$\tilde{x}_i = x_i - \bar{x} + \bar{z}$$

$$\tilde{y}_i = y_i - \bar{y} + \bar{z}$$

Two-samples bootstrap t test (one sided)


Histogram of t.boot



```
> B<-1000
> t.boot<-c(1:B)
> for(b in 1:B)
+ {
+
+ x.b<-sample(x.tilde,m,replace=TRUE)
+ y.b<-sample(y.tilde,n,replace=TRUE)
+ t.boot[b]<-t.test(x.b,y.b)$statistic
+ }
> hist(t.boot,nclass=50,probability=T)
> lines(c(t.obs,t.obs),c(0,1),lwd=3,col=2)
> Pmc<-(1+sum(t.boot > t.obs))/(B+1)
> Pmc
[1] 0.1508492
```

Example 3: Two populations problem: an example of a ratio test

- Population of the most populated city in US in 1920 and 1930.
- Main question: do we see a growth in the population from 1920 to 1930 ?



External dataset:

```
city<-read.table('C:/projects/cim/UpdatesSlides_2017/Data/city.txt',  
header=FALSE,na.strings=".", dec=".", strip.white=TRUE,  
col.names=c("pop1920","pop1930"))
```


Ratio test statistic

$U_i = \text{pop. } 1920$

$X_i = \text{pop. } 1930$

Parameter:

$$\theta = \frac{E(X)}{E(U)} = \frac{\mu_{1930}}{\mu_{1920}}$$


Statistic:

$$t = \frac{\bar{x}}{\bar{u}}$$

```
> t.obs <- mean(x)/mean(u)
```

```
> t.obs
```

```
[1] 1.239019
```



Ratio test statistic


The observed ratio is equal to 1.239019.

Formulation of the hypotheses:

$$H_0 : \frac{\mu_{1920}}{\mu_{1930}} = 1$$

$$H_1 : \frac{\mu_{1920}}{\mu_{1930}} < 1$$

$$t = \frac{\bar{u}}{\bar{x}}$$



Bootstrap algorithm

- Bootstrap algorithm.
- Bootstrap pairs (why ?).


```
> B=5000
> t.boot<-c(1:B)
> for(b in 1:B)
+ {
+   index<-sample(c(1:n),size=n,replace=T)
+   u.boot<-u[index]
+   x.boot<-x[index]
+   t.boot[b]<-mean(x.boot)/mean(u.boot)
+ }
```

Bootstrap p-value

$$P = \frac{1 + \#\{t_b^* > t^{obs}\}}{B + 1}$$

```
> (1+sum(t.boot > t.obs))/(B+1)  
[1] 0.7476505
```

What is wrong here?



City population: paired t-test

$$Z_i = U_i - X_i = \text{pop.}_{1920} - \text{pop.}_{1930}$$

Under the null hypothesis:

$$E(Z_i) = E(U_i - X_i) = 0$$

$$H_0 : \frac{\mu_{1920}}{\mu_{1930}} = 1$$

$$H_1 : \frac{\mu_{1920}}{\mu_{1930}} < 1$$



$$H_0 : \mu_z = 0$$

$$H_1 : \mu_z < 0$$

Each city was measured twice (1920 & 1930).

Under the null hypothesis (no growth) the mean of the difference is expected to be zero.

Resampling under the null

How can we generate the bootstrap sample under the null hypothesis ?

$$H_0 : \mu_Z = 0$$

$$H_0 : \mu_Z > 0$$

$$\tilde{Z}_i = Z_i - \bar{Z} + \mu_0$$

$$E(\tilde{Z}_i) = \mu_0$$

```
> z<-u-x  
> mean(z)  
[1] -24.65306  
> z.tilde<-z-mean(z)  
> mean(z.tilde)  
[1] 5.794596e-16
```

Two samples paired t-test (two sided)


```
> t.test(z,mu=0)
```

One Sample t-test

```
data:  z
t = -6.0276, df = 48, p-value = 2.275e-07
alternative hypothesis: true mean is not equal to 0
95 percent confidence interval:
 -32.87662 -16.42950
sample estimates:
mean of x
-24.65306
```

We reject the null hypothesis

Bootstrap algorithm



```
> Pmc<-(1+sum(abs(t.boot) > abs(t.obs)))/(B+1)
> Pmc
[1] 0.00059988
```

What is the bootstrap algorithm of this problem ?

How can we make sure that we bootstrap under the null hypothesis ?

```
> t.boot<-c(1:B)
> for(b in 1:B)
+ {
+   z.b<-
+     sample(z.tilde,size=nz,replace=T)
+   t.boot[b]<-
+     t.test(z.b,mu=0)$statistic
+ }
>
```


How can we test the original null hypothesis ?

Recall that the original hypotheses were formulated in terms of the ratios.

$$H_0 : \frac{\mu_{1920}}{\mu_{1930}} = 1$$
$$H_1 : \frac{\mu_{1920}}{\mu_{1930}} < 1$$

Test statistic:

$$t = \frac{\bar{x}}{\bar{u}}$$

Bootstrapping under the null ?
Homework 3

Example 3: correlation test

The cars data

The data give the speed of **50** cars and the distances taken to stop.


The data were recorded in the 1920s.

In R:

```
> help(cars)
```

Observed correlation

```
> x<-cars$speed  
> y<-cars$dist  
> cor(x,y)  
[1] 0.8068949  $\rightarrow \hat{\rho}$ 
```




Bootstrap C.I. (non parametric)

```
> quantile(coeff.boot, probs=c(0.025, 0.975))  
      2.5%      97.5%  
0.6906236 0.8831039
```

Non parametric bootstrap
algorithm.

Resampling pairs.



The hypotheses

$$H_0 : \rho = 0$$

$$H_1 : \rho \neq 0$$



Under the null hypothesis the speed and stopping distance are not correlated.

This should be reflected in the bootstrap algorithm.

Non parametric bootstrap algorithm: generate a bootstrap sample under the null hypothesis

The observed sample


x_1 y_1

x_2 y_2

• •

• •

• •

x_{50} y_{50}


We resample x_i and y_i
with replacement

n=50


x_1^* y_1^*

x_2^* y_2^*

• •

• •


• •

x_{50}^* y_{50}^*


The bootstrap sample

The bootstrap algorithm

The bootstrap sample

$$\begin{array}{cc} x_1^* & y_1^* \\ x_2^* & y_2^* \\ \cdot & \cdot \\ \cdot & \cdot \\ \cdot & \cdot \\ x_{50}^* & y_{50}^* \end{array}$$


For each bootstrap sample we calculate the correlation


$$\hat{\rho}_b^*(x^*, y^*)$$

Bootstrap p value (two sided)

$$P = \frac{1 + \#\left\{|\hat{\rho}^*| > |\hat{\rho}^{obs}|\right\}}{B + 1}$$

The bootstrap distribution of the
test statistics **rho**

$\hat{\rho} = 0.806$



R code for the C.I. and bootstrap test: non parametric bootstrap


Bootstrap confidence intervals

```
n<-length(x)
B<-1000
coeff.boot<-c(1:B)
index<-c(1:n)
for(b in 1:B)
{
  index.b<- sample(index,size=n,replace=TRUE)
  x.boot<- x[index.b]
  y.boot<- y[index.b]
  coeff.boot[b]<-cor(x.boot,y.boot)
}
```

Bootstrap test

```
n<-length(x)
B<-1000
coeff.boot<-c(1:B)
index<-c(1:n)
for(b in 1:B)
{
  x.boot<- sample(x,size=n,replace=TRUE)
  y.boot<- sample(y,size=n,replace=TRUE)
  coeff.boot[b]<-cor(x.boot,y.boot)
}
```

Permutation tests



Topics

- Permutation tests for the two samples problems.
- Bootstrap without replacement and permutation tests.

Two populations problem


$$F \rightarrow z = (z_1, z_2, \dots, z_n)$$

$$G \rightarrow y = (y_1, y_2, \dots, y_m)$$

Shift model:


$$H_0 : F = G$$

$$H_1 : F > G$$



The mouse data

```
> z<-c(94,197,16,38,99,141,23)
> y<-c(52,104,146,10,51,30,40,27,46)
> z
[1] 94 197 16 38 99 141 23
> y
[1] 52 104 146 10 51 30 40 27 46
> boxplot(z,y)
```



- 16 mice randomly assigned to treatment and control group.
- Survival time following a test surgery.

Hypotheses and test statistic

$$\begin{array}{l} H_0 : F = G \\ H_1 : F > G \end{array} \quad \longrightarrow \quad \begin{array}{l} H_0 : F(\mu_F) = G(\mu_G) \\ H_1 : F(\mu_F) > G(\mu_G) \end{array} \quad \longrightarrow \quad \begin{array}{l} H_0 : \mu_F = \mu_G \\ H_1 : \mu_F > \mu_G \end{array}$$

$$\hat{\theta} = \bar{z} - \bar{y}$$

$$\hat{\theta} \sim N\left(0, \sigma^2\left(\frac{1}{n} + \frac{1}{m}\right)\right)$$

Standard error is unknown

We use the parameter estimate for the standard error

$$\frac{\hat{\theta}}{se(\hat{\theta})} = \frac{\hat{\theta}}{\hat{\sigma}(\hat{\theta})} \sim N(0,1) \quad \longrightarrow \quad \frac{\hat{\theta}}{\hat{\sigma}(\hat{\theta})} \sim t_{n+m-2}$$

The p value (achieved significance level)


$$ASL = P_{H_0}(\theta \geq \hat{\theta})$$

$$ASL = P_{H_0}\left(t_{14} \geq \frac{\hat{\theta}}{\hat{\sigma}(\hat{\theta})}\right) = P_{H_0}\left(t_{14} \geq \frac{30.63}{54.21\sqrt{\frac{1}{9} + \frac{1}{7}}}\right) = 0.141$$

The p value (achieved significance level)

- The distribution of the test statistic under the null.
- The p value:

$$P_{H_0} \left(t \geq t^{obs} \right) = P \left(\theta \geq \hat{\theta} \right)$$



Two samples problem: one sided test

```
> t.test(z,y,alternative="greater")
```

Welch Two Sample t-test $\rightarrow \sigma_1 \neq \sigma_2$

data: z and y

t = 1.0587, df = 9.654, p-value = 0.1578

alternative hypothesis: true

difference in means is greater than 0


95 percent confidence interval: -22.00168

Inf

sample estimates:

mean of x mean of y

86.85714 56.22222



Permutations

- We combine the two samples.
- Under the null hypothesis the survival time of any mice could come equally from either two groups.

```
> data.frame(sort(x), r.x[order(x)])
```

```
sort.x. r.x.order.x..
```

1	10	y
2	16	z
3	23	z
4	27	y
5	30	y
6	38	z
7	40	y
8	46	y
9	51	y
10	52	y
11	94	z
12	99	z
13	104	y
14	141	z
15	146	y
16	197	z

Permutations

We define two vectors:

Response (survival time)

$$v = (v_1, v_2, \dots, v_N)$$

$$v_i = \begin{cases} y_i & \text{observation} \in G \\ z_i & \text{observation} \in F \end{cases}$$

Membership

$$g = (g_1, g_2, \dots, g_N)$$

$$g_i = \begin{cases} y & \text{observation} \in G \\ z & \text{observation} \in F \end{cases}$$

Permutations

Under $H_0:F=G$, the vector g has equal probability, π , of being equal to one of its possible values.

$$\pi = \frac{1}{\binom{N}{n}} \quad \text{n: sample size from F}$$


Under $H_0:F=G$, each permutation of the data has the same probability (all permutations of z 's and y 's are equally likely)

Test Statistic

$$\hat{\theta} = S(g, v)$$

The test statistic is a function of the two vectors.

The mean difference.

$$\hat{\theta} = \frac{1}{n} \sum_{g_i=z} v_i - \frac{1}{m} \sum_{g_i=y} v_i$$


Permutations

$$\hat{\theta}^* = S(g^*, v)$$

$$ALS = \frac{\#(\hat{\theta}^* \geq \hat{\theta})}{\binom{N}{n}}$$

If we do **all possible permutations** than the p value is the number of time that the test statistic (in the permutation) is grater or equal to the observed test statistic divided by the number of all possible permutations.

What is the problem here ?

Permutations


$$\hat{\theta}^*(b) = S(g^*(b), v)$$

$$ALS = \frac{\#(\hat{\theta}^*(b) \geq \hat{\theta})}{B}$$

We draw B samples **without replacement** (why ?) from the vector v .

The first n are assigned to the first group the last m are assigned to the second group.

Permutations



```
> x<-c(z,y)
> B<-1000
> t.boot<-c(1:B)
> for(b in 1:B)
+ {
+   x.boot<-sample(x,size=n+m,replace=FALSE)
+   z.boot<-x.boot[1:n]
+   y.boot<-x.boot[(n+1):(n+m)]
+   t.boot[b]<-
+     t.test(z.boot,y.boot,alternative="greater",
+           var.equal=TRUE)$statistic
+ }
> hist(t.boot,nclass=50,probability=TRUE)
> lines(c(t.obs,t.obs),c(0,10),col=3,lwd=3)
```

```
> (1+sum(t.boot>t.obs))/(B+1)
[1] 0.1348651
```


Extra example: binary data

Two populations and two proportions

Population 1

$$X_i = \begin{cases} 1 & \pi_1 \\ 0 & 1 - \pi_1 \end{cases}$$



Sample 1

$$X_1, X_2, \dots, X_{n_1}$$

Population 2

$$Y_i = \begin{cases} 1 & \pi_2 \\ 0 & 1 - \pi_2 \end{cases}$$



Sample 2

$$Y_1, Y_2, \dots, Y_{n_2}$$

The testing problem

The null hypothesis that we want to test

$$H_0 : \pi_2 - \pi_1 = 0$$

$$H_1 : \pi_2 - \pi_1 \neq 0$$

The sample proportions

- Suppose S_1 the number of successes in the first sample and S_2 the number of successes in the second sample.
- The proportions of the sample are then given by

$$\hat{\pi}_1 = \bar{P}_1 = \frac{1}{n} \sum_{i=1}^n x_i = \frac{\text{the number of "1s" in sample}_1}{n_1}$$

$$\hat{\pi}_1 = \bar{P}_1 = \frac{S_1}{n_1}$$

$$\hat{\pi}_2 = \bar{P}_2 = \frac{S_2}{n_2}$$

The variance of the samples means

$$X_1, X_2, \dots, X_{n_1}$$

a sample from the first population

$$E(\bar{P}_1) = \pi_1$$

$$Var(\bar{P}_1) = \frac{\pi_1(1 - \pi_1)}{n_1}$$

$$Y_1, Y_2, \dots, Y_{n_2}$$

a sample from the second population

$$E(\bar{P}_2) = \pi_2$$

$$Var(\bar{P}_2) = \frac{\pi_2(1 - \pi_2)}{n_2}$$

The expected value and the variance of the difference

$$E(\bar{P}_2 - \bar{P}_1) = \pi_2 - \pi_1$$

$$Var(\bar{P}_1 - \bar{P}_2) = Var(\bar{P}_2) + Var(\bar{P}_1) = \frac{\pi_2(1 - \pi_2)}{n_2} + \frac{\pi_1(1 - \pi_1)}{n_1}$$

The proportion under H_0

The null hypothesis says that $\pi_1 = \pi_2 = \pi$ (for some unknown value π).

Under H_0 :

$$E(\bar{P}_2 - \bar{P}_1) = 0$$

$$Var(\bar{P}_2 - \bar{P}_1) = Var(\bar{P}_2) + Var(\bar{P}_1) = \pi(1 - \pi) \left(\frac{1}{n_1} + \frac{1}{n_2} \right)$$

$$H_0 : \pi_2 - \pi_1 = 0 \Rightarrow \pi_2 - \pi_1 = \pi$$


The pooled sample proportion

Problem is that π is not known. An estimator for π is given by the **pooled sample proportion**

$$\bar{P} = \frac{n_1 \bar{P}_1 + n_2 \bar{P}_2}{n_1 + n_2} = \frac{S_1 + S_2}{n_1 + n_2}$$

The test statistic

$$\frac{\bar{P}_2 - \bar{P}_1 - (p_2 - p_1)_{H_0}}{\sqrt{\pi(1-\pi)\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}} = \frac{\bar{P}_2 - \bar{P}_1}{\sqrt{\pi(1-\pi)\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}} \approx N(0,1)$$



unknown parameter



$$\frac{\bar{P}_2 - \bar{P}_1}{\sqrt{\bar{P}(1-\bar{P})\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}} \sim N(0,1)$$

Example:
Agresti – An Introduction to
Categorical Data Analysis
Aspirin use and heart attacks.

Example: aspirin use and myocardial infarction (MI)

Table 2.3. Cross Classification of Aspirin Use and Myocardial Infarction

Group	Myocardial Infarction		Total
	Yes	No	
Placebo	189	10,845	11,034
Aspirin	104	10,933	11,037

Source: Preliminary Report: Findings from the Aspirin Component of the Ongoing Physicians' Health Study. *New Engl. J. Med.*, 318: 262-264, 1988.

Relationship between aspirin use and myocardial infarction (heart attacks).

A five-year randomized study testing whether regular intake of aspirin reduces mortality from cardiovascular disease.

Population 1

$$X_i = \begin{cases} 1 & \pi_1 \\ 0 & 1 - \pi_1 \end{cases}$$

Population 2

$$Y_i = \begin{cases} 1 & \pi_2 \\ 0 & 1 - \pi_2 \end{cases}$$

Every other day, the male physicians participating in the study took either one aspirin tablet or a placebo.

$$\pi_\ell = P(\text{heart attack}), \quad \ell = \text{Placebo, Aspirin}$$

Example: aspirin use and heart attacks.

Table 2.3. Cross Classification of Aspirin Use and Myocardial Infarction

Group	Myocardial Infarction		Total
	Yes	No	
Placebo	189	10,845	11,034
Aspirin	104	10,933	11,037

Source: Preliminary Report: Findings from the Aspirin Component of the Ongoing Physicians' Health Study. *New Engl. J. Med.*, 318: 262-264, 1988.

Main interest:

$$H_0 : \pi_{Placebo} - \pi_{Aspirin} = 0$$

```

> P.Placebo<-189/(189+10845)
> P.Placebo
[1] 0.01712887
> P.Aspirin<-104/(104+10933)
> P.Aspirin
[1] 0.00942285
> P.Placebo-P.Aspirin
[1] 0.007706024

```

$\hat{\pi}_1 = \hat{P}(\text{heart attack} \mid \text{Placebo})$
 $\hat{\pi}_2 = \hat{P}(\text{heart attack} \mid \text{Aspirin})$
 $\hat{\pi}_1 - \hat{\pi}_2$

Proportion under the null and the test statistic

Table 2.3. Cross Classification of Aspirin Use and Myocardial Infarction

Group	Myocardial Infarction		Total
	Yes	No	
Placebo	189	10,845	11,034
Aspirin	104	10,933	11,037

Source: Preliminary Report: Findings from the Aspirin Component of the Ongoing Physicians' Health Study. *New Engl. J. Med.*, 318: 262-264, 1988.

```
> P.tot<- (189+104) / (189+10845+104+10933)
> P.tot
[1] 0.01327534
```

$$\bar{P} = \frac{n_1 \bar{P}_1 + n_2 \bar{P}_2}{n_1 + n_2} = \frac{S_1 + S_2}{n_1 + n_2}$$

```
> z<- (P.Placebo-P.Aspirin)/sqrt(P.tot*(1-P.tot)*(1/n.P+1/n.A))
> z
[1] 5.001388
```


$$z = \frac{\bar{P}_2 - \bar{P}_1}{\sqrt{\bar{P}(1-\bar{P})\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}$$

```
> 1-pnorm(z,0,1)
[1] 2.845948e-07
```

p-value

The null hypothesis is rejected.


Parametric bootstrap test: populations and samples




Parametric bootstrap test

Population 1

$$X_i = \begin{cases} 1 & \pi \\ 0 & 1 - \pi \end{cases}$$



$$\begin{aligned} H_0 : \pi_2 - \pi_1 &= 0 \\ H_1 : \pi_2 - \pi_1 &\neq 0 \end{aligned}$$



Population 2

$$Y_i = \begin{cases} 1 & \pi \\ 0 & 1 - \pi \end{cases}$$

$$\hat{\pi} = \bar{P} = \frac{n_1 \bar{P}_1 + n_2 \bar{P}_2}{n_1 + n_2} = \frac{S_1 + S_2}{n_1 + n_2}$$

Table 2.3. Cross Classification of Aspirin Use and Myocardial Infarction

Group	Myocardial Infarction		Total
	Yes	No	
Placebo	189	10,845	11,034
Aspirin	104	10,933	11,037

Source: Preliminary Report: Findings from the Aspirin Component of the Ongoing Physicians' Health Study. *New Engl. J. Med.*, 318: 262-264, 1988.

$$n_1 = 11034$$

$$n_1 = 11037$$

$$\hat{\pi} = \bar{P} = \frac{189 + 104}{11034 + 11037}$$

Parametric bootstrap test

For 1:B

Population 1

$$x_b^* \sim B(11034, \hat{\pi})$$

$$\hat{p}_{1b}^* = \frac{1}{n_1} \sum x$$

$$\begin{aligned} H_0 : \pi_2 - \pi_1 &= 0 \\ H_1 : \pi_2 - \pi_1 &\neq 0 \end{aligned}$$

Population 2

$$y_b^* \sim B(11037, \hat{\pi})$$

$$\hat{p}_{2b}^* = \frac{1}{n_2} \sum y$$

$$p_b^* = \frac{S_{1b}^* + S_{2b}^*}{n_1 + n_2}$$

$$z_b^* \frac{\hat{p}_{1b}^* - \hat{p}_{2b}^*}{\sqrt{p_b^*(1-p_b^*)\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}$$

Test statistic for the
b'th bootstrap
sample.

The bootstrap loop


Two sided parametric bootstrap test

Simulate B samples and calculate the test statistic values:

$$z_1^*, z_2^*, \dots, z_B^*$$

Bootstrap p value (for two sided test):

$$P = \frac{1 + \# \left\{ |z_b^*| > |z^{obs}| \right\}}{B + 1}$$



R code: parametric bootstrap

```

B<-10000
z.b<-c(1:B)
for(i in 1:B)
{
  S1<-sum(rbinom(n.P,1,P.tot))
  P.P.b<-S1/(189+10845)
  S2<-sum(rbinom(n.A,1,P.tot))
  P.A.b<-S2/(104+10933)
  P.tot.b<-(S1+S2)/(n.P+n.A)
  z.b[i]<-(P.P.b-P.A.b)/sqrt(P.tot.b*(1-P.tot.b)*(1/n.P+1/n.A))
}

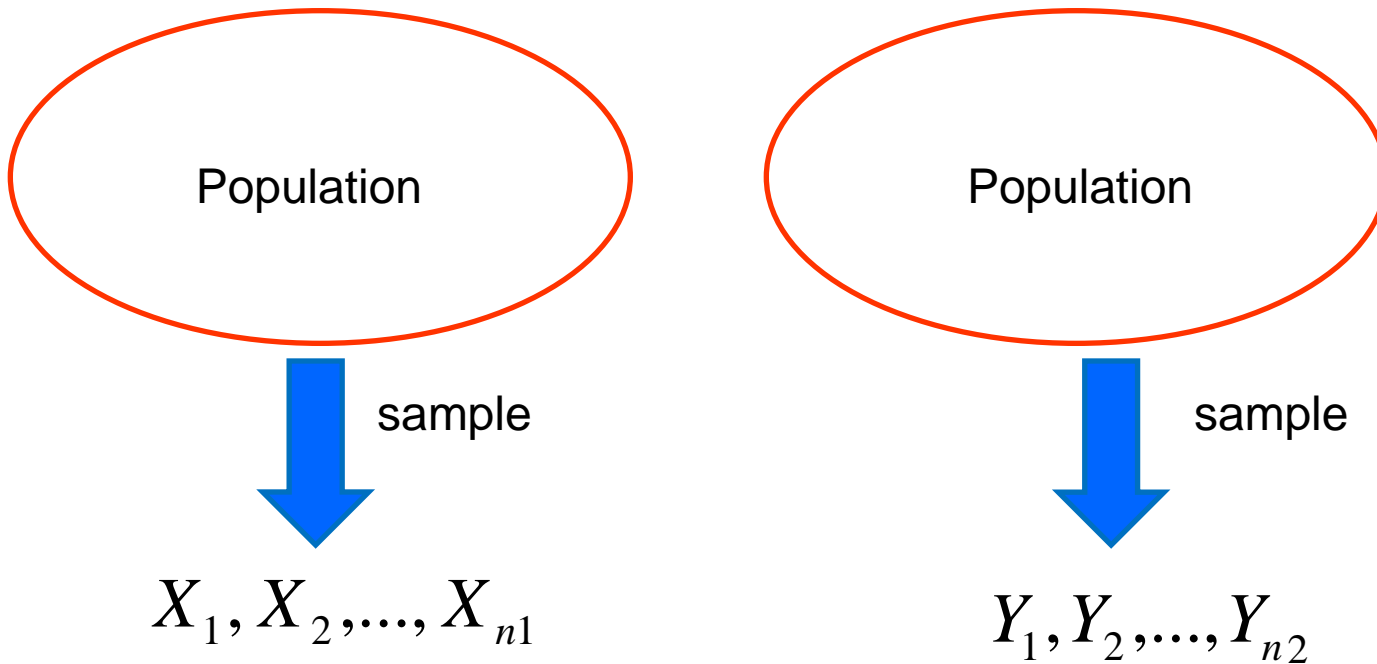
```

$$x_b^* \sim B(11034, \hat{\pi})$$

$$y_b^* \sim B(11037, \hat{\pi})$$

$$z_b^* = \frac{\hat{p}_{1b}^* - \hat{p}_{2b}^*}{\sqrt{p_b^*(1-p_b^*)\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}$$

Non parametric bootstrap: populations and samples



$$(Z_1, Z_2, \dots, Z_{n1}, Z_{n1+1}, \dots, Z_{n1+n2}) = (X_1, X_2, \dots, X_{n1}, Y_1, \dots, Y_{n2})$$

Application of algorithm 16.1

$$z = (y, x)$$

Simulate B samples of size $n_1 + n_2$ with replacement from z.

The first n_1 are considered as observations x^* and the remaining n_2 as observations y^* .

Calculate the test statistic for each bootstrap sample:

$$t(z^*)_1, t(z^*)_2, \dots, t(z^*)_B$$

Test statistic for the b'th bootstrap sample:

$$t(z^*)_b = z_b^* \frac{\hat{p}_{1b}^* - \hat{p}_{2b}^*}{\sqrt{p_b^*(1 - p_b^*) \left(\frac{1}{n_1} + \frac{1}{n_2} \right)}}$$


Two sided non parametric bootstrap test

Simulate B samples and calculate the test statistic values:

$$z_1^*, z_2^*, \dots, z_B^*$$

Bootstrap p value (for two sided test):


$$P = \frac{1 + \# \left\{ |z_b^*| > |z^{obs}| \right\}}{B + 1}$$




The observed test statistic is larger than all bootstrap statistics

⇒ $P = \frac{1}{B + 1}$

R code for non parametric bootstrap

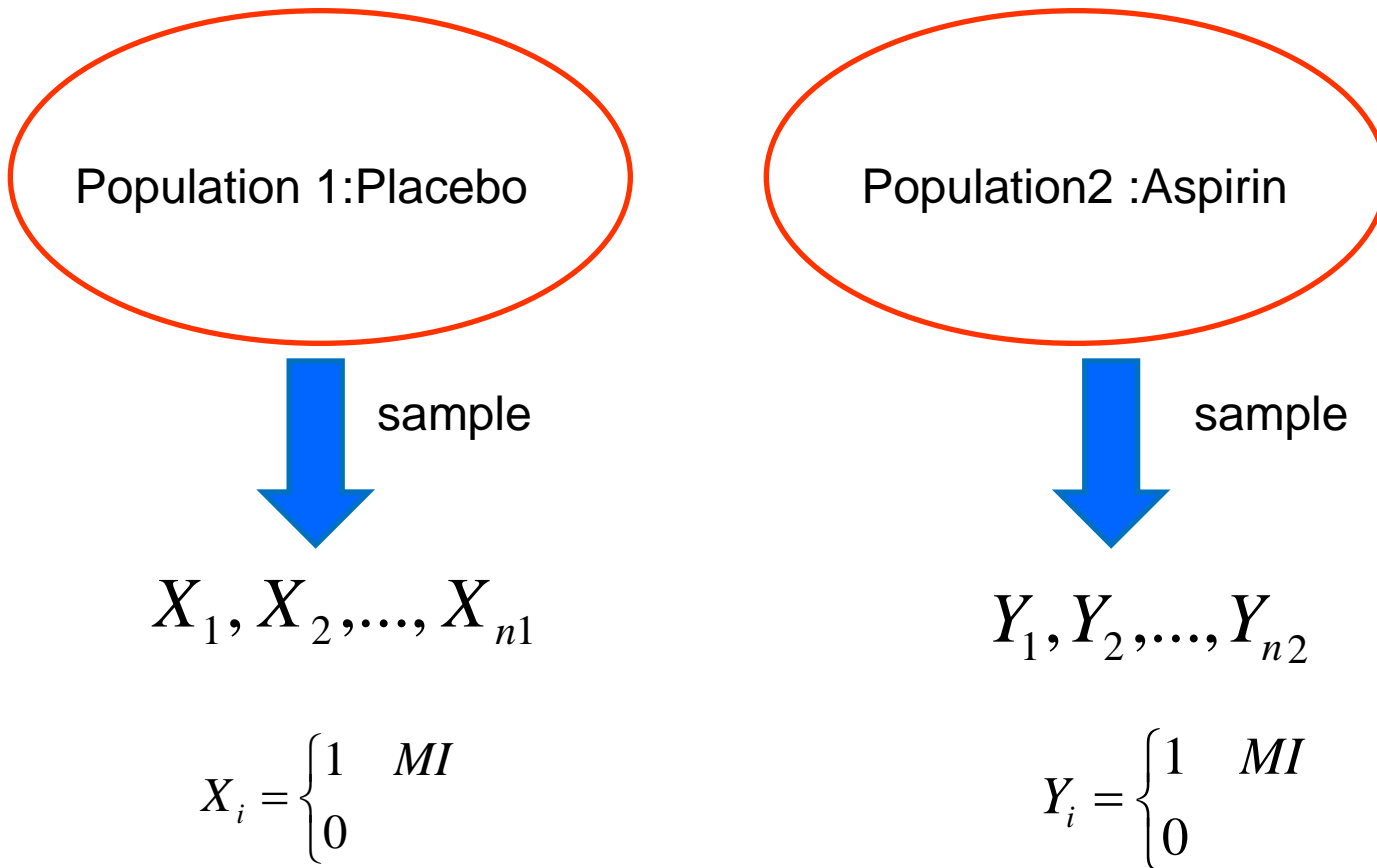
```
z.i<-c(Placebo,Aspirin)      (Z1,Z2,...,Zn1,Zn1+1,...,Zn1+n2)=(X1,X2,...,Xn1,Y1,...,Yn2)
nz<-length(z.i)
nz  n1 + n2
```

```
B<-10000
z.b<-c(1:B)
for(i in 1:B)
{
  z.boot<-sample(z.i,size=nz,replace =TRUE)
  x.boot<-z.boot[1:n.P]
  y.boot<-z.boot[(n.P+1):nz]
  S1<-sum(x.boot)
  P.P.b<-S1/(189+10845)
  S2<-sum(y.boot)
  P.A.b<-S2/(104+10933)
  P.tot.b<-(S1+S2)/(n.P+n.A)
  z.b[i]<-(P.P.b-P.A.b)/sqrt(P.tot.b*(1-P.tot.b)*(1/n.P+1/n.A))
}
```

$$z_b^* \frac{\hat{p}_{1b}^* - \hat{p}_{2b}^*}{\sqrt{p_b^*(1-p_b^*)\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}$$


```
hist(z.b,nclass=50,xlim=c(-6,6),probability=TRUE)
lines(c(z,z),c(0,100),col=2,lwd=2)
```

Permutations test: populations and samples



$$v = (X_1, X_2, \dots, X_{n1}, Y_1, \dots, Y_{n2})$$

Permutations

We define two vectors:

Response (MI)

$$v = (v_1, v_2, \dots, v_N)$$

$$v_i = \begin{cases} x_i & \text{observation} \in P \\ y_i & \text{observation} \in A \end{cases}$$

$$v_i = \begin{cases} 1 & MI \\ 0 & \text{not } MI \end{cases}$$

Membership

$$g = (g_1, g_2, \dots, g_N)$$

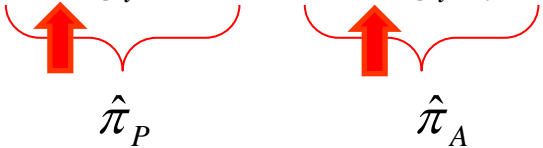
$$g_i = \begin{cases} x & \text{observation} \in P \\ y & \text{observation} \in A \end{cases}$$

Test Statistic

$$\hat{\theta} = S(g, v)$$

The test statistic is a function of the two vectors.

The mean difference $(\hat{\pi}_P - \hat{\pi}_A)$.

$$\hat{\theta} = \frac{1}{n_1} \sum_{g_i=x} v_i - \frac{1}{n_2} \sum_{g_i=y} v_i$$


The diagram illustrates the relationship between the sample sizes n_1 and n_2 and the estimated proportions $\hat{\pi}_P$ and $\hat{\pi}_A$. Red arrows point from $\hat{\pi}_P$ to n_1 and from $\hat{\pi}_A$ to n_2 . Red curly braces are placed under n_1 and n_2 , indicating that these sample sizes are used to calculate the estimated proportions $\hat{\pi}_P$ and $\hat{\pi}_A$ respectively.

Permutations

$$\hat{\theta}^*(b) = S(g^*(b), v)$$

We draw B samples **without replacement** (why ?) from the vector v .

$$z_b^* = \hat{\theta}^*(b) = \frac{\hat{p}_{1b}^* - \hat{p}_{2b}^*}{\sqrt{p_b^*(1-p_b^*)\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}$$

The first n are assigned to the first group the last m are assigned to the second group.

$$ALS = \frac{\#\left(\hat{\theta}^*(b) \geq \hat{\theta}\right)}{B}$$


Two sided non parametric bootstrap test

Simulate B samples and calculate the test statistic values:


$$z_1^*, z_2^*, \dots, z_B^*$$

Bootstrap p value (for two sided test):

$$P = \frac{1 + \# \left\{ |z_b^*| > |z^{obs}| \right\}}{B + 1}$$



R code for non parametric bootstrap

```
z.i<-c(Placebo,Aspirin)  v=(X1,X2,...,Xn1,Y1,...,Yn2)
nz<-length(z.i)
nz  n1+n2
```

```
for(i in 1:B)
{
```

```
z.boot<-sample(z.i,size=nz,replace =FALSE)
```

```
x.boot<-z.boot[1:n.P]
```

```
y.boot<-z.boot[(n.P+1):nz]
```

```
S1<-sum(x.boot)
```

```
P.P.b<-S1/(189+10845)
```

```
S2<-sum(y.boot)
```

```
P.A.b<-S2/(104+10933)
```

```
P.tot.b<-(S1+S2)/(n.P+n.A)
```

```
z.b[i]<-(P.P.b-P.A.b)/sqrt(P.tot.b*(1-P.tot.b)*(1/n.P+1/n.A))
```

```
}
```

$$z_b^* \frac{\hat{p}_{1b}^* - \hat{p}_{2b}^*}{\sqrt{p_b^*(1-p_b^*)\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}$$