

Introduction to R: the tidyverse package

Ziv Shkedy, Hasselt University (July 2020)

```
## Warning: package 'mvtnorm' was built under R version 3.6.2
## Warning: package 'tidyverse' was built under R version 3.6.3
## -- Attaching packages ----- tidyverse 1.3.0 --
## v tibble  2.1.3      v dplyr    1.0.0
## v tidyr   1.1.0      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0
## v purrr   0.3.4
## Warning: package 'tidyr' was built under R version 3.6.3
## Warning: package 'purrr' was built under R version 3.6.3
## Warning: package 'dplyr' was built under R version 3.6.3
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
## Warning: package 'dslabs' was built under R version 3.6.3
## Warning: package 'NHANES' was built under R version 3.6.3
```

Introduction

The tidyverse

The tidyverse is a collection of R packages that work in data frames in a tidy format. All packages are uploaded in CRAN, and can be installed using `install.packages()`. In this chapter of the interactive book, we cover materials at an introduction level and follow closely the topics presented in Chapter 4 in the book *Data Analysis and Prediction Algorithms with R* by Rafael A. Irizarry.

What do we cover in this chapter ?

The chapter is developed at a beginner level, We cover few functions from the tidyverse packages and illustrate the basic concepts using different examples of the following functions:

- `mutate()`
- `filter()`
- `select()`
- The pipe: `%>%`
- `summarize()`
- `group_by()`
- `arrange()`
- `top_n()`

Our aim in this tutorial is not to teach `ggplot2` (this will be done in a different chapter). However, some functions of the package are used to visualize the main patterns in the datasets we used to illustrate the examples presented in this chapter. The following graphical functions are used for visualization:

- `qplot()`
- `ggplot() + geom_jitter()`
- `ggplot() + geom_point()`
- `ggplot() + geom_density()`
- `stripplot()`

Online references

Materials about tidyverse are widely available online. We list below a selection that we find useful and clear.

YouTube tutorials: tidyverse in R

- For a YouTube tutorial about tidyverse in R by Mark Gingrass see [YTtidyverse1](#).
- For a YouTube tutorial about tidyverse in R by Garreth Golemund see [YTtidyverse2a](#).
- For a YouTube tutorial about tidyverse in R by Garreth Golemund see [YTtidyverse2b](#).
- For a YouTube tutorial about tidyverse in R by Ben Stenhaus see [YTtidyverse3](#).

Online book

Chapter 4 in the book: *Data Analysis and Prediction Algorithms with R* by Rafael A. Irizarry see [Booktidyverse1](#).

Datasets

Many datasets are used for illustration in this chapter. All of them are data frames available in R. We do not focus in this chapter on the question how to read the data but rather on the question how to organize the data for the analysis and basic visualization techniques.

The murders data

The murders dataset gives information about the number of gun murders in 51 US states (2010).

```
data("murders")
head(murders)
```

```
##      state abb region population total
## 1  Alabama AL  South    4779736    135
## 2  Alaska  AK   West     710231     19
## 3  Arizona AZ   West    6392017    232
## 4  Arkansas AR  South    2915918     93
## 5 California CA  West    37253956   1257
## 6  Colorado CO   West    5029196     65
```

In total, five variables are included in the data.

```
dim(murders)
```

```
## [1] 51  5
```

The heights data

The heights dataset gives information about the self reported heights (in inches) for males and females of 1050 subjects.

```
data(heights)
dim(heights)
```

```
## [1] 1050  2
```

the first 6 subjects are shown below.

```
head(heights)
```

```
##      sex height
## 1  Male      75
## 2  Male      70
## 3  Male      68
## 4  Male      74
## 5  Male      61
## 6 Female      65
```

The NHANES data

The NHANES dataset consists of data from the US National Health and Nutrition Examination Study. Information about 76 variables is available for 10000 subjects included in the study.

```
library(NHANES)
data(NHANES)
dim(NHANES)
```

```
## [1] 10000    76
```

Variables names are listed below.

```
names(NHANES)
```

```
## [1] "ID"          "SurveyYr"      "Gender"        "Age"
## [5] "AgeDecade"    "AgeMonths"     "Race1"         "Race3"
## [9] "Education"    "MaritalStatus" "HHIncome"      "HHIncomeMid"
## [13] "Poverty"      "HomeRooms"     "HomeOwn"       "Work"
## [17] "Weight"       "Length"        "HeadCirc"      "Height"
## [21] "BMI"          "BMICatUnder20yrs" "BMI_WHO"       "Pulse"
## [25] "BPSysAve"     "BPDiaAve"      "BPSys1"        "BPDia1"
## [29] "BPSys2"       "BPDia2"        "BPSys3"        "BPDia3"
## [33] "Testosterone" "DirectChol"    "TotChol"       "UrineVol1"
## [37] "UrineFlow1"   "UrineVol2"     "UrineFlow2"    "Diabetes"
## [41] "DiabetesAge"  "HealthGen"     "DaysPhysHlthBad" "DaysMentHlthBad"
## [45] "LittleInterest" "Depressed"     "nPregnancies"  "nBabies"
## [49] "Age1stBaby"   "SleepHrsNight" "SleepTrouble"  "PhysActive"
## [53] "PhysActiveDays" "TVHrsDay"      "CompHrsDay"    "TVHrsDayChild"
## [57] "CompHrsDayChild" "Alcohol12PlusYr" "AlcoholDay"    "AlcoholYear"
## [61] "SmokeNow"     "Smoke100"      "Smoke100n"     "SmokeAge"
## [65] "Marijuana"    "AgeFirstMarij" "RegularMarij"  "AgeRegMarij"
## [69] "HardDrugs"    "SexEver"       "SexAge"        "SexNumPartnLife"
## [73] "SexNumPartYear" "SameSex"       "SexOrientation" "PregnantNow"
```

The Chicks Weights data

The Chick Weights data is a data frame in R called `chickwts`. The data consists of 71 newly hatched chicks that were randomly allocated into six groups. Each group was given a different feed supplement. Their weights (the response variable) in grams after six weeks are given along with feed types (the factor).

```
head(chickwts)
```

```
##   weight      feed
## 1    179 horsebean
## 2    160 horsebean
## 3    136 horsebean
## 4    227 horsebean
## 5    217 horsebean
## 6    168 horsebean
```

```
dim(chickwts)
```

```
## [1] 71  2
```

The Chicken Weights data

The `ChickWeight` dataset (the R object `ChickWeight`) is a data frame with 578 rows and 4 columns from an experiment on the effect of diet on early growth of chicks. In total, 50 subjects were measured 12 times over a period of 21 days.

```
head(ChickWeight)
```

```
## Grouped Data: weight ~ Time | Chick
##   weight Time Chick Diet
## 1     42    0     1    1
## 2     51    2     1    1
## 3     59    4     1    1
## 4     64    6     1    1
## 5     76    8     1    1
## 6     93   10     1    1
```

Note that each row in the data represents the chick weight in a specific day.

```
dim(ChickWeight)
```

```
## [1] 578  4
```

The cars data

The cars data (the R object `mtcars`) gives information about fuel consumption (`mpg`) and 10 aspects of automobile design and performance for 32 automobiles (1973-74 models).

```
dim(mtcars)
```

```
## [1] 32 11
```

```
head(mtcars)
```

##		mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
##	Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
##	Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
##	Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
##	Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
##	Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
##	Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

Data manipulation with the tidyverse package

Tidy data

Tidy data is a data format in which each row represents one measurement for one observation and columns are, as usual, the variables in the data.

The murders data

The murders is an example of a tidy data. The information for each state is given in one line. Note that in this case, each observation (=state) has information in one data line.

```
data("murders")
head(murders)
```

```
##      state abb region population total
## 1  Alabama AL  South    4779736    135
## 2  Alaska  AK   West     710231     19
## 3  Arizona AZ   West    6392017    232
## 4  Arkansas AR  South    2915918     93
## 5 California CA   West    37253956  1257
## 6  Colorado CO   West     5029196     65
```

The murder rate by region is shown in the stripplot presented in Figure~@ref(fig:fig1b) that shows clearly that in the west, the murder rate is the lowest.

```
ggplot(murders, aes(region,population)) + geom_jitter(position = position_jitter(width = .05))
```

The ChickWeight data

In the Chicken Weight data, each observation is a chick and it was measured in 12 times points.

```
unique(ChickWeight$Time)
```

```
## [1] 0 2 4 6 8 10 12 14 16 18 20 21
```

In the data frame, each measurement is presented in one data line. Hence, the ChickWeight is a tidy data. This implies that the data for each observation is presented in 12 lines (note that NOT ALL subjects were measured in 12 times). Data for the first 6 time points of the first chick are listed below.

```
head(ChickWeight)
```

```
## Grouped Data: weight ~ Time | Chick
##   weight Time Chick Diet
## 1    42    0     1    1
## 2    51    2     1    1
## 3    59    4     1    1
## 4    64    6     1    1
## 5    76    8     1    1
## 6    93   10     1    1
```

The boxplot of the chicken weights by time point, presented in Figure~@ref(fig:fig1c), shows the increasing trend of the weight over time.

```
ggplot(ChickWeight, aes(as.factor(Time),weight)) + geom_boxplot()
```




Figure 1: Dotplot using the gg2plot package




Figure 2: Boxplot for the singers data (I.1).

Adding a variable (column)

Suppose that we would like to calculate the murder rate per 100000 people that is

$$\frac{\text{total}}{\text{population}} \times 100000.$$

This can be done using the `mutate()` function that has the general call of: `mutate(data frame , new variable)` .

The murders data

For the murders data we have

```
data("murders")
murders <- mutate(murders, rate = total / population * 100000)
```

Note that after calculating the murder rate, the murders has an extra column (=variable) for the rate.

```
head(murders)

##      state abb region population total    rate
## 1  Alabama  AL  South   4779736   135 2.824424
## 2  Alaska   AK   West    710231    19 2.675186
## 3  Arizona  AZ   West   6392017   232 3.629527
## 4  Arkansas AR  South   2915918    93 3.189390
## 5 California CA  West  37253956  1257 3.374138
## 6  Colorado CO   West   5029196    65 1.292453
```

The NHANES data

The BMI of a person is given by

$$BMI = \frac{\text{weight}}{\text{height}^2}.$$

To calculate the BMI in the NHANES we use

```
data("NHANES")
Data_new <- mutate(NHANES, BMI_new = Weight / (Height*Height))
```

The histogram of the BMI is shown in Figure~@ref(fig:fig1d).

```
qplot(BMI_new , data=Data_new, geom="histogram")
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
## Warning: Removed 366 rows containing non-finite values (stat_bin).
```

Practical session

For the cars data (`mtcars`):

- Define a new variable: mile per gallon / weight.
- Produce a boxplot for the variable mile per gallon / weight by number of number of cylinders.

For the murders data :




Figure 3: Histogram for the BMI.

- Produce a boxplot for the murder rate by region.

Filtering

Selection of observations from the data can be done using the function `filter()` .

The murders data

For murder dataset, we can select all the states with murder rate ≤ 0.71 by

```
filter(murders, rate <= 0.71)
```

##	state	abb	region	population	total	rate
## 1	Hawaii	HI	West	1360301	7	0.5145920
## 2	Iowa	IA	North Central	3046355	21	0.6893484
## 3	New Hampshire	NH	Northeast	1316470	5	0.3798036
## 4	North Dakota	ND	North Central	672591	4	0.5947151
## 5	Vermont	VT	Northeast	625741	2	0.3196211

The cars data

For cars data, suppose that we would like to plot the cars' weight versus the cars mpg for cars with weight ≤ 3 . We first select all cars with weight smaller or equal to 3:

```
mtcars1<-filter(mtcars, wt <= 3)
```

The new data frame mtcars1 contains the information for all cars with weight lower than 3,

```
mtcars1
```

##		mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
##	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
##	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
##	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
##	Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
##	Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
##	Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
##	Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
##	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
##	Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
##	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
##	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
##	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

The scatterplot in Figure~@ref(fig:fig1e) below of the weight versus the mpg can be produce using the following code.

```
ggplot(mtcars1, aes(x=wt, y=mpg)) +
  geom_point( color="#69b3a2")
```

Practical session

For the NHANES data:

- Select all the observations for which the BMI is greater or equal to 30.




Figure 4: mile per galllon versus weight.

- Produce the histogram for the variable BMI for all male.

For the chick weight data (chickwts):

- Select all chicks with weight smaller than 158.
- Select all chicks from the horsebean feed group.

Selecting columns

In the previous section we use the function `filter()` to select observations. In this section we focus on variable selection from the data frame using the function `select()`.

The murders data

Originally, the murder data frame has 6 variables.

```
dim(murders)
```

```
## [1] 51 6
```

We define a new data frame `newdata` that contains only 3 of the variables in the murder data frame.

```
newdata <- select(murders, state, region, rate)
dim(newdata)
```

```
## [1] 51 3
```

Note that we can further filter the observations, for example a selection of all observations with murder rate lower than 0.71:

```
filter(newdata, rate <= 0.71)
```

```
##           state      region    rate
## 1      Hawaii      West 0.5145920
## 2      Iowa North Central 0.6893484
## 3 New Hampshire Northeast 0.3798036
## 4 North Dakota North Central 0.5947151
## 5      Vermont Northeast 0.3196211
```

The NHANES data

In this example, we define a new data frame with contains 6 variables from the NHANES .

```
NHANES1 <- select(NHANES, Gender, Age, Weight, Height, BMI, Diabetes)
dim(NHANES1)
```

```
## [1] 10000 6
```

```
head(NHANES1)
```

```
## # A tibble: 6 x 6
##   Gender Age Weight Height BMI Diabetes
##   <fct> <int> <dbl> <dbl> <dbl> <fct>
## 1 male   34   87.4   165.   32.2 No
## 2 male   34   87.4   165.   32.2 No
## 3 male   34   87.4   165.   32.2 No
## 4 male    4   17    105.   15.3 No
## 5 female 49   86.7   168.   30.6 No
```

```
## 6 male      9   29.8   133.  16.8 No
```

A density plot of the BMI by gender is shown in Figure~@ref(fig:fig1fa).

```
ggplot(data=NHANES1, aes(x=BMI, group=Gender, fill=Gender)) +  
  geom_density(adjust=1.5)
```

```
## Warning: Removed 366 rows containing non-finite values (stat_density).
```




Figure 5: Density plot for the BMI by gender (I).

Alternatively, we can present the density in a separate panel per gender group as shown in Figure~@ref(fig:fig1fb).

```
ggplot(data=NHANES1, aes(x=BMI, group=Gender, fill=Gender)) +  
  geom_density(adjust=1.5)+  
  facet_wrap(~Gender)
```

```
## Warning: Removed 366 rows containing non-finite values (stat_density).
```

Practical session

For the The cars dataset:

- Create a new data frame which contains the variables: Miles/(US) gallon, Number of cylinders and Displacement (cu.in.).

For the NHANES data:




Figure 6: Density plot for the BMI by gender (II).

- Create a new data frame in which you include the gender and two other continuous variables that you choose.
- For each of the continuous variable that you choose, produce density plot by gender group.

The pipe: %>%

In the previous section we use the functions `filter()` and `select()` to select a part of the dataset in two steps. In this section we use the pipe `%>%` to make the selection in one step.

The murders data

We want to select from the murder data frame three variables and all the states with murder rate lower or equal to 0.71. The two selection conditions can be applied to the data frame in one step in the following way:

```
murders %>% select(state, region, rate) %>% filter(rate <= 0.71)
```

```
##           state      region    rate
## 1      Hawaii      West 0.5145920
## 2      Iowa North Central 0.6893484
## 3 New Hampshire Northeast 0.3798036
## 4 North Dakota North Central 0.5947151
## 5      Vermont Northeast 0.3196211
```

The region is a factor variable with four levels

```
murders$region
```

```
## [1] South      West      West      South      West
## [6] West      Northeast South      South      South
## [11] South      West      West      North Central North Central
## [16] North Central North Central South      South      Northeast
## [21] South      Northeast North Central North Central South
## [26] North Central West      North Central West      Northeast
## [31] Northeast West      Northeast South      North Central
## [36] North Central South      West      Northeast Northeast
## [41] South      North Central South      South      West
## [46] Northeast South      West      South      North Central
## [51] West
## Levels: Northeast South North Central West
```

We define a new data frame which contains the states from the Northeast and West regions

```
data1<-filter(murders, region %in% c("Northeast", "West"))
data1
```

```
##           state abb  region population total    rate
## 1      Alaska AK    West    710231    19 2.6751860
## 2      Arizona AZ    West    6392017   232 3.6295273
## 3      California CA    West   37253956  1257 3.3741383
## 4      Colorado CO    West    5029196    65 1.2924531
## 5      Connecticut CT Northeast   3574097    97 2.7139722
## 6      Hawaii HI     West    1360301     7 0.5145920
## 7      Idaho ID     West    1567582    12 0.7655102
## 8      Maine ME     Northeast   1328361    11 0.8280881
## 9      Massachusetts MA Northeast   6547629   118 1.8021791
## 10     Montana MT     West     989415    12 1.2128379
```



```
## 11      Nevada NV      West      2700551      84 3.1104763
## 12 New Hampshire NH Northeast 1316470      5 0.3798036
## 13      New Jersey NJ Northeast 8791894     246 2.7980319
## 14      New Mexico NM      West      2059179      67 3.2537239
## 15      New York NY Northeast 19378102    517 2.6679599
## 16      Oregon OR      West      3831074      36 0.9396843
## 17 Pennsylvania PA Northeast 12702379    457 3.5977513
## 18 Rhode Island RI Northeast 1052567      16 1.5200933
## 19      Utah UT      West      2763885      22 0.7959810
## 20      Vermont VT Northeast 625741       2 0.3196211
## 21      Washington WA      West      6724540     93 1.3829942
## 22      Wyoming WY      West      563626       5 0.8871131
```

We select all states from the Northeast and West regions with mtder rate lower or equal to 1

```
filter(data1, rate <=1)
```

```
##      state abb      region population total      rate
## 1      Hawaii HI      West      1360301       7 0.5145920
## 2      Idaho ID      West      1567582      12 0.7655102
## 3      Maine ME Northeast 1328361      11 0.8280881
## 4 New Hampshire NH Northeast 1316470       5 0.3798036
## 5      Oregon OR      West      3831074      36 0.9396843
## 6      Utah UT      West      2763885      22 0.7959810
## 7      Vermont VT Northeast 625741       2 0.3196211
## 8      Wyoming WY      West      563626       5 0.8871131
```

In one step, the selection above can be implemented with the following code:

```
data2<-filter(murders, region %in% c("Northeast", "West") & rate <= 1)
print(data2)
```

```
##      state abb      region population total      rate
## 1      Hawaii HI      West      1360301       7 0.5145920
## 2      Idaho ID      West      1567582      12 0.7655102
## 3      Maine ME Northeast 1328361      11 0.8280881
## 4 New Hampshire NH Northeast 1316470       5 0.3798036
## 5      Oregon OR      West      3831074      36 0.9396843
## 6      Utah UT      West      2763885      22 0.7959810
## 7      Vermont VT Northeast 625741       2 0.3196211
## 8      Wyoming WY      West      563626       5 0.8871131
```

```
select(data2,state,region,population)
```

```
##      state      region population
## 1      Hawaii      West      1360301
## 2      Idaho      West      1567582
## 3      Maine Northeast 1328361
## 4 New Hampshire Northeast 1316470
## 5      Oregon      West      3831074
## 6      Utah      West      2763885
## 7      Vermont Northeast 625741
## 8      Wyoming      West      563626
```

The NHANES data

We select 6 variables for all female in the NHANES data frame

```
NHANES1<-NHANES %>% select(Gender, Age, Weight, Height,BMI,Diabetes) %>% filter(Gender %in% c("female"))
head(NHANES1)
```

```
## # A tibble: 6 x 6
##   Gender   Age Weight Height   BMI Diabetes
##   <fct> <int> <dbl> <dbl> <dbl> <fct>
## 1 female    49  86.7  168.  30.6 No
## 2 female    45  75.7  167.  27.2 No
## 3 female    45  75.7  167.  27.2 No
## 4 female    45  75.7  167.  27.2 No
## 5 female    10  38.6  142.  19.2 No
## 6 female    58  57.5  148.  26.2 No
```

Figure~@ref(fig:fig1g) shows the distribution of the BMI by diabetes group for the feample in the NHANES data frame.

```
stripplot(Diabetes ~ jitter(BMI),
           data = NHANES1 ,
           aspect = 1, jitter = T,
           xlab="BMI", col = 1)
```




Figure 7: BMI by Diabetes group (for female).

Practical session

For the The ChickWeight dataset (the R object ChickWeight):

- Create a new data frame which contains all measurement from the first 12 days of the first 20 subjects.
- For the new data frame, produce a histogram for the chick's weight by day.

For the cars data (mtcars):

- Select all cars with more than 3 forward gears and weight smaller than 3.

The summarize function

The function `summarize()` allows us to produce summary statistics for variables in the data frame.

heights data

The height data frame gives the height by gender for 1050 individuals

```
library(dplyr)
library(dslabs)
data(heights)
dim(heights)
```

```
## [1] 1050    2
```

The first 6 observations in the data are listed below.

```
head(heights)
```

```
##      sex height
## 1  Male     75
## 2  Male     70
## 3  Male     68
## 4  Male     74
## 5  Male     61
## 6 Female    65
```

We can calculate the mean and standard deviation for female using the function `summarize()`. Note that we first filter the data and define a subgroup contains the data for female

```
s <- heights %>% filter(sex == "Female") %>%
  summarize(average = mean(height), standard_deviation = sd(height))
```

The object `s` stores the results

```
s
```

```
##      average standard_deviation
## 1 64.93942      3.760656
```

We can define a vector that contains the results

```
c(s$average, s$standard_deviation)
```

```
## [1] 64.939424 3.760656
```

Alternatively, we can define a vector with the female heights (`height.female`) and calculate the mean and standard deviation for this vector.

```
height.female <- heights$height[heights$sex == "Female"]
mean(height.female)
```

```
## [1] 64.93942
sd(height.female)
```

```
## [1] 3.760656
```

The median, minimum and maximum height for female

```
heights %>%
  filter(sex == "Female") %>% summarize(median = median(height), minimum = min(height),
                                       maximum = max(height))
```

```
##      median minimum maximum
## 1 64.98031      51      79
```

These summary statistics can be also calculate can the function quantile .

```
heights %>% filter(sex == "Female") %>%
  summarize(range = quantile(height, c(0, 0.5, 1)))
```

```
##      range
## 1 51.00000
## 2 64.98031
## 3 79.00000
```

The chicks data

To calculate the mean and standard deviation for the chick weights we use

```
s <- chickwts %>% summarize(average = mean(weight), standard_deviation = sd(weight))
s
```

```
##      average standard_deviation
## 1 261.3099      78.0737
```

Note that for this example we ignore the diet group.

Practical session

For the NHANES data:

- Calculate the mean, meadian and standard deviation of the variable BMI.
- Calculate the mean, meadian and standard deviation of the variable BMI only for male.

For the cars data:

- Calculate the mean, meadian and standard deviation of the variable weight.
- Calculate the mean, meadian and standard deviation of the variable weight only for cars with automatic transmission.

Analysis by group

In this section we focus on an analysis in which the analysis is done across a level of a factor in the data frame. For example, the diet group in the chick data frame etc.

The heights data

The mean and standard deviation for the height by gender can be calculate by adding the function `group_by(sex)`

```
heights %>%
  group_by(sex) %>%
  summarize(average = mean(height), standard_deviation = sd(height))

## `summarise()` ungrouping output (override with `.groups` argument)

## # A tibble: 2 x 3
##   sex      average standard_deviation
##   <fct>      <dbl>          <dbl>
## 1 Female    64.9            3.76
## 2 Male     69.3            3.61
```

The same results can be obtained using the function `tapply` .

```
tapply(heights$height,heights$sex,mean)

##   Female      Male
## 64.93942 69.31475

tapply(heights$height,heights$sex,sd)

##   Female      Male
## 3.760656 3.611024
```

The murders data

The median murder rate by region using the `group_by(region)` and the `summarize()` functions

```
murders %>% group_by(region) %>%
  summarize(median_rate = median(rate))

## `summarise()` ungrouping output (override with `.groups` argument)

## # A tibble: 4 x 2
##   region      median_rate
##   <fct>          <dbl>
## 1 Northeast      1.80
## 2 South          3.40
## 3 North Central  1.97
## 4 West           1.29
```

The median murder rate by region using the `tapply()` function.

```
tapply(murders$rate,murders$region,median)

##      Northeast      South North Central      West
##      1.802179      3.398069      1.971105      1.292453
```

The chicks data

Summary statistics by diet group

```
chickwts %>%
  group_by(feed) %>%
  summarize(average = mean(weight), standard_deviation = sd(weight))

## `summarise()` ungrouping output (override with `.groups` argument)

## # A tibble: 6 x 3
##   feed      average standard_deviation
##   <fct>      <dbl>          <dbl>
## 1 casein     324.            64.4
## 2 horsebean  160.            38.6
## 3 linseed    219.            52.2
## 4 meatmeal   277.            64.9
## 5 soybean    246.            54.1
## 6 sunflower  329.            48.8
```

The stripplot in Figure~@ref(fig:figh) reveals that the weights in the horsebean group are, in general, the smallest in the sample.

```
stripplot(feed ~ jitter(weight),
  data = chickwts,
  aspect = 1, jitter = T,
  xlab="Chicks weight", col = 1)
```




Figure 8: by diet group.

Practical session

For the NHANES data:

- Calculate the mean BMI by gender.
- Calculate the mean age by diabetes status (the variable Diabetes).

For the cars data:

- Calculate the mean, median and standard deviation of the variable weight by transmission type (automatic vs. manual).
- Calculate the mean, median and standard deviation of the variable miles per gallon by number of cylinders.

Sorting

The murders data

We can sort a data frame by a variable x using the function `arrange(x)`. For the murder data frame, we sort the data by the population size

```
murders %>%  
  arrange(population) %>% head()
```

##	state	abb	region	population	total	rate
## 1	Wyoming	WY	West	563626	5	0.8871131
## 2	District of Columbia	DC	South	601723	99	16.4527532
## 3	Vermont	VT	Northeast	625741	2	0.3196211
## 4	North Dakota	ND	North Central	672591	4	0.5947151
## 5	Alaska	AK	West	710231	19	2.6751860
## 6	South Dakota	SD	North Central	814180	8	0.9825837

The same sorting can be implemented using the function `order()`. In this case the rows will be presented in the order of the population.

```
data1<-murders[order(murders$population),]  
head(data1)
```

##	state	abb	region	population	total	rate
## 51	Wyoming	WY	West	563626	5	0.8871131
## 9	District of Columbia	DC	South	601723	99	16.4527532
## 46	Vermont	VT	Northeast	625741	2	0.3196211
## 35	North Dakota	ND	North Central	672591	4	0.5947151
## 2	Alaska	AK	West	710231	19	2.6751860
## 42	South Dakota	SD	North Central	814180	8	0.9825837

We sort the data frame by rate from the lowest to the highest rate

```
murders %>%  
  arrange(rate) %>%  
  head()
```

##	state	abb	region	population	total	rate
## 1	Vermont	VT	Northeast	625741	2	0.3196211
## 2	New Hampshire	NH	Northeast	1316470	5	0.3798036
## 3	Hawaii	HI	West	1360301	7	0.5145920
## 4	North Dakota	ND	North Central	672591	4	0.5947151
## 5	Iowa	IA	North Central	3046355	21	0.6893484

```
## 6          Idaho  ID          West    1567582    12 0.7655102
```

We can change the order using the function `desc()` so the data are presented from the highest to the lowest rate in a decreasing order.

```
murders %>%
  arrange(desc(rate)) %>%
  head
```

```
##           state abb      region population total      rate
## 1 District of Columbia DC      South    601723    99 16.452753
## 2      Louisiana LA      South    4533372   351  7.742581
## 3      Missouri MO North Central    5988927   321  5.359892
## 4      Maryland MD      South    5773552   293  5.074866
## 5 South Carolina SC      South    4625364   207  4.475323
## 6      Delaware DE      South    897934    38  4.231937
```

The Chicken Weight data

The first 6 observations in the chicken weight data belongs to the first chick at time point 0 to 10.

```
head(ChickWeight)
```

```
## Grouped Data: weight ~ Time | Chick
##   weight Time Chick Diet
## 1     42    0     1    1
## 2     51    2     1    1
## 3     59    4     1    1
## 4     64    6     1    1
## 5     76    8     1    1
## 6     93   10     1    1
```

We sort the data frame according to the Time variable. After sorting, the first 6 lines in the data frame are the measurements for chock 1-6 at day 21.

```
ChickWeight %>%
  arrange(desc(Time)) %>%
  head
```

```
## Grouped Data: weight ~ Time | Chick
##   weight Time Chick Diet
## 1    205   21     1    1
## 2    215   21     2    1
## 3    202   21     3    1
## 4    157   21     4    1
## 5    223   21     5    1
## 6    157   21     6    1
```

We can reverse the order, in this case the first 6 lines are the measurements for check 1-6 at baseline (Time =0).

```
ChickWeight %>%
  arrange(Time) %>%
  head
```

```
## Grouped Data: weight ~ Time | Chick
##   weight Time Chick Diet
## 1     42    0     1    1
```



```
## 2      40      0      2      1
## 3      43      0      3      1
## 4      42      0      4      1
## 5      41      0      5      1
## 6      41      0      6      1
```

Practical session

- Sort the NHANES data according to the subjects' BMI.
- For the NHANES data, select all subjects with diabetes (the variable Diabetes) and sort according to the subjects' age.
- Sort the cars data according to the cars' mpg.

Nested sorting

The murders data

Suppose that we want to present the data in an increasing order of x across a level of a factor y . We can sort the data frame by a variable x within the factor levels using the function `arrange(y,x)`. For the murder data frame, we sort the data by murder rate within the region

```
murders %>%
  arrange(region, rate) %>%
  head()
```

```
##           state abb   region population total      rate
## 1      Vermont  VT Northeast    625741      2 0.3196211
## 2 New Hampshire NH Northeast    1316470      5 0.3798036
## 3         Maine ME Northeast    1328361     11 0.8280881
## 4  Rhode Island RI Northeast    1052567     16 1.5200933
## 5 Massachusetts MA Northeast    6547629    118 1.8021791
## 6      New York NY Northeast    19378102   517 2.6679599
```

The cars data

Figure~@ref(fig:figi) shows that mpg as the number of cylinders decreases.

```
stripplot(cyl ~ jitter(mpg),
  data = mtcars,
  aspect = 1, jitter = T,
  xlab="MPG", col = 1)
```

We can sort the cars according to their mpg (in an increasing order) by the number of cylinders

```
mtcars %>%
  arrange(cyl, mpg)
```

```
##           mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## Volvo 142E    21.4   4 121.0 109 4.11 2.780 18.60 1  1   4    2
## Toyota Corona 21.5   4 120.1  97 3.70 2.465 20.01 1  0   3    1
## Datsun 710    22.8   4 108.0  93 3.85 2.320 18.61 1  1   4    1
## Merc 230      22.8   4 140.8  95 3.92 3.150 22.90 1  0   4    2
## Merc 240D     24.4   4 146.7  62 3.69 3.190 20.00 1  0   4    2
## Porsche 914-2 26.0   4 120.3  91 4.43 2.140 16.70 0  1   5    2
```




Figure 9: MPG by number of cylinders.

## Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
## Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
## Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
## Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
## Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
## Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
## Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
## Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
## Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
## Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
## Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
## Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
## Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
## Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
## Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
## Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
## Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
## Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
## Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
## AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
## Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
## Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
## Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
## Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
## Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
## Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2

Practical session

- Sort the NHANES data according to the subjects' BMI by Gender.
- Sort the NHANES data according to the subjects' age by smoking status (the variable Smoke100).
- Sort the Chicks Weights data (chickwts) according to the chicks' weight by diet group.

Top n

The murders data

To print the top n observations according to the variable x we can use the function `top_n(n, x)`. For the murders data, we print the top 5 states with the highest murder rate

```
murders %>% top_n(5, rate)
```

##	state	abb	region	population	total	rate
## 1	District of Columbia	DC	South	601723	99	16.452753
## 2	Louisiana	LA	South	4533372	351	7.742581
## 3	Maryland	MD	South	5773552	293	5.074866
## 4	Missouri	MO	North Central	5988927	321	5.359892
## 5	South Carolina	SC	South	4625364	207	4.475323

The cars data

Figure~@ref(fig:figj) shows the scatterplot of the cars' weight versus the cars' mpg.

```
ggplot(mtcars, aes(x=wt, y=mpg)) +  
  geom_point()
```




Figure 10: MPG by number of cylinders.

The top 4 cars, with the highest mpg are given below

```
mtcars %>% top_n(4,mpg)
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
## Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
## Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
## Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2

Practical session

- Print the 10 observations in the NHANES data with the highest BMI.
- Create a new data frame that contains the top 5 observations with the older age in in the NHANES data.